


# tf.talk()

An Introduction to Deep Learning  
with TensorFlow




# Table of Contents

# Table of Catents




Theory

# Table of Catents



Theory



Practice

# Background

# Background

- ▶ CS Student @ TUM

# Background

- ▶ CS Student @ TUM
- ▶ Google & Bloomberg Intern

# Background

- ▶ CS Student @ TUM
- ▶ Google & Bloomberg Intern

Seminar Topic: *Deep Learning With TensorFlow*

`github.com/peter-can-write/tensorflow-paper`




# Neural Networks


# Neural Networks




# Neural Networks




# Neural Networks




# Neural Networks




# Neural Networks



# Neural Networks




# Neural Networks



$$\begin{bmatrix} x_0 & x_1 \end{bmatrix} \underset{\mathbf{x}}{\times} \begin{bmatrix} w_{0,0} & w_{0,1} \\ w_{1,0} & w_{1,1} \end{bmatrix} \underset{\mathbf{W}}{+} \begin{bmatrix} b_0 \\ b_1 \end{bmatrix} \underset{\mathbf{b}}{=} \begin{bmatrix} y_0 \\ y_1 \end{bmatrix} \underset{\mathbf{y}}{=}$$




# Neural Networks




$$\begin{matrix} \begin{bmatrix} x_0 & x_1 \end{bmatrix} \\ \mathbf{x} \end{matrix} \times \begin{matrix} \begin{bmatrix} w_{0,0} & w_{0,1} \\ w_{1,0} & w_{1,1} \end{bmatrix} \\ \mathbf{W} \end{matrix} + \begin{matrix} \begin{bmatrix} b_0 \\ b_1 \end{bmatrix} \\ \mathbf{b} \end{matrix} = \begin{matrix} \begin{bmatrix} y_0 \\ y_1 \end{bmatrix} \\ \mathbf{y} \end{matrix} \rightsquigarrow \hat{\mathbf{y}}$$


# Neural Networks




# Neural Networks




# Neural Networks




# Neural Networks



# Deep Neural Networks



# Deep Neural Networks



# Deep Neural Networks


Deep Learning assumes that data is structured



Deep Learning assumes that data is structured

hierarchically

# Deep Neural Networks




# Deep Neural Networks

# Deep Neural Networks




# Deep Neural Networks




# Deep Neural Networks

74	28	75	10	84	58	17	0	60	41	26	10	4	51	21	56
94	70	53	4	7	45	75	23	99	17	97	86	68	59	83	25
68	52	41	27	27	26	9	93	61	50	29	50	22	2	90	32
9	61	49	80	97	4	66	45	59	50	7	30	25	25	76	51
10	89	48	85	41	12	12	86	48	80	55	20	79	46	73	61
99	26	84	66	45	12	18	66	33	18	46	82	9	61	95	97
75	86	55	17	70	92	37	72	95	41	18	91	99	99	23	46
59	50	33	51	49	77	11	81	98	76	30	29	74	53	13	17
57	40	4	25	83	49	66	38	33	49	9	90	55	88	41	79
88	6	83	49	29	53	33	11	67	60	31	32	41	47	1	50
66	44	92	72	67	61	79	43	56	35	92	1	82	17	65	29
34	23	78	56	63	86	47	68	77	22	2	69	36	27	78	65
37	84	46	41	98	49	82	32	17	77	38	66	13	57	46	4
7	44	29	21	98	71	39	52	23	32	16	77	87	35	57	32
25	6	84	55	7	90	74	71	4	27	76	67	34	32	0	14
47	41	75	72	98	75	57	19	94	26	92	97	41	75	1	36

# Deep Neural Networks




# Deep Neural Networks






# Deep Neural Networks



# Deep Neural Networks



# Deep Neural Networks

- ▶ We want to classify images into one of  $k$  classes

# Deep Neural Networks

- ▶ We want to classify images into one of  $k$  classes
- ▶ Extract hierarchical features


# Deep Neural Networks

- ▶ We want to classify images into one of  $k$  classes
- ▶ Extract hierarchical features

Edges


# Deep Neural Networks

- ▶ We want to classify images into one of  $k$  classes
- ▶ Extract hierarchical features




# Deep Neural Networks

- ▶ We want to classify images into one of  $k$  classes
- ▶ Extract hierarchical features



# Deep Neural Networks

- ▶ We want to classify images into one of  $k$  classes
- ▶ Extract hierarchical features





Why reinvent the wheel?

Why reinvent the wheel?

40	54	64
38	97	19
99	41	59


# Deep Neural Networks

Why reinvent the wheel?

12
74
3
61
28
81
49
47
39


# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?



# Deep Neural Networks


Why reinvent the wheel?






# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?



# Deep Neural Networks


Why reinvent the wheel?






# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?



# Deep Neural Networks

Why reinvent the wheel?




Time for ze kitteh



This is a cat ♥

Time for ze kitteh



Still a cat ♥♥

Time for ze kitteh



Half cat / half salad ♥♥♥♥

Time for ze kitteh




Many cats ♥♥♥♥♥

## Weight Sharing




# Convolutional Neural Networks

## Weight Sharing




# Convolutional Neural Networks

## Weight Sharing




# Convolutional Neural Networks

## Weight Sharing




# Convolutional Neural Networks

## Weight Sharing




# Convolutional Neural Networks

## Weight Sharing



# Convolutional Neural Networks: Mechanics



Image

# Convolutional Neural Networks: Mechanics

0.4	0.9	0.1
0.7	0.2	0.6
0.8	0.3	0.5

Image

# Convolutional Neural Networks: Mechanics

0.4	0.9	0.1
0.7	0.2	0.6
0.8	0.3	0.5

Image

5.7	2.4
3.1	0.9

Kernel



# Convolutional Neural Networks: Mechanics

$5.7 \cdot 0.4$	$2.4 \cdot 0.9$	0.1
$3.1 \cdot 0.7$	$0.9 \cdot 0.2$	0.6
0.8	0.3	0.5

Image

# Convolutional Neural Networks: Mechanics

$5.7 \cdot 0.4$	$2.4 \cdot 0.9$	0.1
$3.1 \cdot 0.7$	$0.9 \cdot 0.2$	0.6
0.8	0.3	0.5

Image

6.79

Output

# Convolutional Neural Networks: Mechanics

0.4	$5.7 \cdot 0.9$	$2.4 \cdot 0.1$
0.7	$3.1 \cdot 0.2$	$0.9 \cdot 0.6$
0.8	0.3	0.5

Image

6.79	6.53
------	------

Output

# Convolutional Neural Networks: Mechanics

0.4	0.9	0.1
$5.7 \cdot 0.7$	$2.4 \cdot 0.2$	0.6
$3.1 \cdot 0.8$	$0.9 \cdot 0.3$	0.5

Image

6.79	6.53
7.67	

Output

# Convolutional Neural Networks: Mechanics


0.4	0.9	0.1
0.7	$5.7 \cdot 0.2$	$2.4 \cdot 0.6$
0.8	$3.1 \cdot 0.3$	$0.9 \cdot 0.5$

Image


6.79	6.53
7.67	3.96

Output


# Convolutional Neural Networks: Mechanics



# Convolutional Neural Networks: Mechanics




# Convolutional Neural Networks: Mechanics







# Convolutional Neural Networks: Mechanics



# Convolutional Neural Networks: Mechanics




# Convolutional Neural Networks: Mechanics







# Convolutional Neural Networks: Mechanics




# Convolutional Neural Networks: Mechanics



# Convolutional Neural Networks: Mechanics




# Convolutional Neural Networks: Mechanics







# Convolutional Neural Networks: Mechanics



# Convolutional Neural Networks: Mechanics



# Convolutional Neural Networks: Mechanics

## Recipe for a Convolutional Layer

# Convolutional Neural Networks: Mechanics

## Recipe for a Convolutional Layer

- ▶ Ingredients

# Convolutional Neural Networks: Mechanics

## Recipe for a Convolutional Layer

- ▶ Ingredients

1. Image  $I$  with dimensions  $w \times h \times d$

## Recipe for a Convolutional Layer

► Ingredients

1. Image  $I$  with dimensions  $w \times h \times d$
2. A kernel (filter)  $K$  of size  $k \times l \times m$

## Recipe for a Convolutional Layer

- ▶ Ingredients

1. Image  $I$  with dimensions  $w \times h \times d$
2. A kernel (filter)  $K$  of size  $k \times l \times m$

- ▶ Cooking

## Recipe for a Convolutional Layer

- ▶ Ingredients
  1. Image  $I$  with dimensions  $w \times h \times d$
  2. A kernel (filter)  $K$  of size  $k \times l \times m$
- ▶ Cooking
  - ▶ Put the image into the oven at  $150^{\circ}\text{C}$



## Recipe for a Convolutional Layer

- ▶ Ingredients
  1. Image  $I$  with dimensions  $w \times h \times d$
  2. A kernel (filter)  $K$  of size  $k \times l \times m$
- ▶ Cooking
  - ▶ Don't put the image into the oven at 150°C

## Recipe for a Convolutional Layer

- ▶ Ingredients
  1. Image  $I$  with dimensions  $w \times h \times d$
  2. A kernel (filter)  $K$  of size  $k \times l \times m$
- ▶ Cooking
  - ▶ Don't put the image into the oven at  $150^{\circ}\text{C}$
  - ▶ Slide the kernel across the image

## Recipe for a Convolutional Layer

- ▶ Ingredients

1. Image  $I$  with dimensions  $w \times h \times d$
2. A kernel (filter)  $K$  of size  $k \times l \times m$

- ▶ Cooking

- ▶ Don't put the image into the oven at  $150^{\circ}\text{C}$
- ▶ Slide the kernel across the image
- ▶ Compute the “dot product” for each configuration

# Convolutional Neural Networks: Pooling

# Convolutional Neural Networks: Pooling

66	2
6	32

# Convolutional Neural Networks: Pooling

66	2
6	32

# Convolutional Neural Networks: Pooling

50	17
66	2

# Convolutional Neural Networks: Pooling

9	50
33	66



# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

66
----

# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

66	79
----	----

# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

66	79
66	

# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

66	79
66	128

# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

66	79
66	128

- ▶ *Pooling* achieves translational invariance

# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

66	79
66	128

- ▶ *Pooling* achieves translational invariance
- ▶ A form of downsampling



# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

66	79
66	128

- ▶ *Pooling* achieves translational invariance
- ▶ A form of downsampling
- ▶ Other pooling functions possible

# Convolutional Neural Networks: Architecture

INPUT  $\rightarrow$  [CONV+  $\rightarrow$  POOL?]+  $\rightarrow$  FC+  $\rightarrow$  OUTPUT



TensorFlow



- ▶ An open source deep learning library

# TensorFlow




- ▶ An open source deep learning library
- ▶ Released by Google in November 2015



- ▶ An open source deep learning library
- ▶ Released by Google in November 2015
- ▶ Especially suited to:  
“Large-scale machine learning on heterogeneous distributed systems”

# Computational Paradigms

# Computational Paradigms




**Computational Graphs**

$$\hat{y} = \tanh(\mathbf{x}^\top \mathbf{w} + b)$$



# Computational Paradigms




$$\hat{y} = \tanh(\mathbf{x}^\top \mathbf{w} + b)$$

## Computational Graphs

### 1. Operations

# Computational Paradigms




$$\hat{y} = \tanh(\mathbf{x}^\top \mathbf{w} + b)$$

## Computational Graphs

1. Operations
2. Tensors

# Computational Paradigms




$$\hat{y} = \tanh(\mathbf{x}^\top \mathbf{w} + b)$$

## Computational Graphs

1. Operations
2. Tensors
3. Variables

# Computational Paradigms



## Computational Graphs

1. Operations
2. Tensors
3. Variables
4. Sessions


$$\hat{y} = \text{session.run}(\tanh(\mathbf{x}^T \mathbf{w} + b))$$

# Execution Model

# Execution Model

**Actors**


# Execution Model



## Actors

1. Client

# Execution Model




## Actors

1. Client
2. Master



# Execution Model




## Actors

1. Client

2. Master

3. Workers


# Execution Model



## Actors

1. Client
2. Master
3. Workers
4. Devices


# Execution Model



## Actors

1. Client
2. Master
3. Workers
4. Devices

# Execution Model



1. Client
2. Master
3. Workers
4. Devices

# Visualization Tools

# Visualization Tools

- ▶ Deep Neural Networks have the tendency of being . . . deep

## Visualization Tools

- ▶ Deep Neural Networks have the tendency of being . . . deep
- ▶ Easy to drown in the complexity of an architecture

# Visualization Tools


- ▶ Deep Neural Networks have the tendency of being . . . deep
- ▶ Easy to drown in the complexity of an architecture
- ▶ > 36,000 nodes for Google's *Inception* model





# Visualization Tools


- ▶ Deep Neural Networks have the tendency of being . . . deep
- ▶ Easy to drown in the complexity of an architecture
- ▶ > 36,000 nodes for Google's *Inception* model














# TensorBoard to the Rescue


# Walkthrough



# Walkthrough


INPUT  $\rightarrow$  [CONV  $\rightarrow$  POOL] $\{2\}$   $\rightarrow$  FC  $\rightarrow$  OUTPUT

# Walkthrough



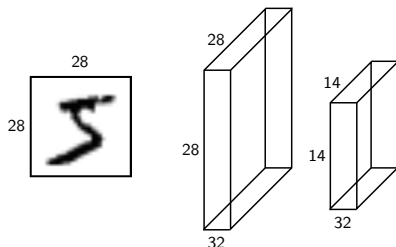
INPUT -> [CONV -> POOL]{2} -> FC -> OUTPUT

# Walkthrough




INPUT  $\rightarrow$  [CONV  $\rightarrow$  POOL] {2}  $\rightarrow$  FC  $\rightarrow$  OUTPUT

# Walkthrough




INPUT -> [CONV -> POOL]{2} -> FC -> OUTPUT

# Walkthrough



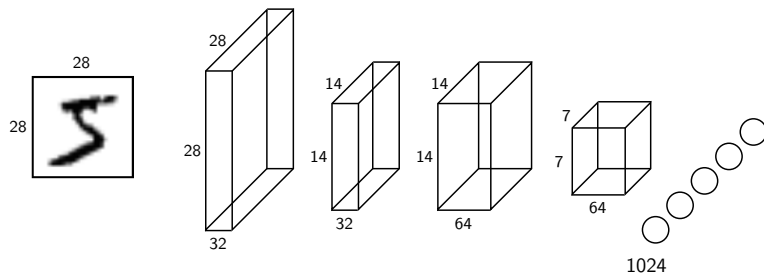
INPUT  $\rightarrow$  [CONV  $\rightarrow$  POOL] $\{2\}$   $\rightarrow$  FC  $\rightarrow$  OUTPUT

# Walkthrough




INPUT  $\rightarrow$  [CONV  $\rightarrow$  POOL] {2}  $\rightarrow$  FC  $\rightarrow$  OUTPUT

# Walkthrough



INPUT  $\rightarrow$  [CONV  $\rightarrow$  POOL]{2}  $\rightarrow$  FC  $\rightarrow$  OUTPUT

# Walkthrough



INPUT  $\rightarrow$  [CONV  $\rightarrow$  POOL] {2}  $\rightarrow$  FC  $\rightarrow$  OUTPUT



How do I continue?

# Resources

# Resources

- ▶ MOOCs
  - ▶ Machine Learning by Andrew Ng @ Coursera
  - ▶ Deep Learning by Google @ Udacity
  - ▶ Machine Learning Nanodegree @ Udacity

# Resources

## ▶ MOOCs

- ▶ Machine Learning by Andrew Ng @ Coursera
- ▶ Deep Learning by Google @ Udacity
- ▶ Machine Learning Nanodegree @ Udacity

## ▶ Websites

- ▶ <http://colah.github.io>
- ▶ <http://cs231n.github.io>
- ▶ <http://karpathy.github.io>
- ▶ <http://www.deeplearningbook.org>
- ▶ <https://www.kaggle.com>
- ▶ <https://www.tensorflow.org>

# Stay in Touch!

- ▶ `peter@goldsborough.me`
- ▶ `linkedin.com/in/petergoldsborough`
- ▶ `github.com/goldsborough`
- ▶ `@peterawks`

# Stay in Touch!

- ▶ `peter@goldsborough.me`
- ▶ `linkedin.com/in/petergoldsborough`
- ▶ `github.com/goldsborough`
- ▶ `@peterawks`

`github.com/peter-can-talk/cambridge-2016`

# Q & A