


# tf.talk()

An Introduction to Deep Learning  
with TensorFlow




# Table of Contents

# Table of Catents




Theory

# Table of Catents



Theory



Practice

# Background

# Background

- ▶ CS Student @ TUM

# Background

- ▶ CS Student @ TUM
- ▶ Google & Bloomberg Intern

# Background

- ▶ CS Student @ TUM
- ▶ Google & Bloomberg Intern

Seminar Topic: *Deep Learning With TensorFlow*

`github.com/peter-can-write/tensorflow-paper`




# Neural Networks


# Neural Networks




# Neural Networks




# Neural Networks




# Neural Networks




# Neural Networks



# Neural Networks




# Neural Networks



$$\begin{bmatrix} x_0 & x_1 \end{bmatrix} \underset{\mathbf{x}}{\times} \begin{bmatrix} w_{0,0} & w_{0,1} \\ w_{1,0} & w_{1,1} \end{bmatrix} \underset{\mathbf{W}}{+} \begin{bmatrix} b_0 \\ b_1 \end{bmatrix} \underset{\mathbf{b}}{=} \begin{bmatrix} y_0 \\ y_1 \end{bmatrix} \underset{\mathbf{y}}{=}$$




# Neural Networks




$$\begin{matrix} \begin{bmatrix} x_0 & x_1 \end{bmatrix} \\ \mathbf{x} \end{matrix} \times \begin{matrix} \begin{bmatrix} w_{0,0} & w_{0,1} \\ w_{1,0} & w_{1,1} \end{bmatrix} \\ \mathbf{W} \end{matrix} + \begin{matrix} \begin{bmatrix} b_0 \\ b_1 \end{bmatrix} \\ \mathbf{b} \end{matrix} = \begin{matrix} \begin{bmatrix} y_0 \\ y_1 \end{bmatrix} \\ \mathbf{y} \end{matrix} \rightsquigarrow \hat{\mathbf{y}}$$


# Neural Networks




# Neural Networks




# Neural Networks




# Neural Networks



# Deep Neural Networks



# Deep Neural Networks



# Deep Neural Networks


Deep Learning assumes that data is structured



Deep Learning assumes that data is structured

hierarchically

# Deep Neural Networks




# Deep Neural Networks

# Deep Neural Networks




# Deep Neural Networks




# Deep Neural Networks

92	73	73	81	57	59	39	33	3	61	48	58	36	70	39	75
38	77	41	91	51	11	8	49	31	43	70	53	0	7	92	85
39	90	53	58	26	51	5	50	14	80	57	5	79	10	81	93
91	31	40	5	37	19	76	27	18	69	1	40	81	38	61	6
9	71	1	19	49	74	36	58	65	28	5	38	83	82	45	44
25	56	65	73	95	39	5	49	94	83	66	6	21	92	62	93
34	22	63	40	96	62	97	50	10	92	21	41	28	92	20	20
55	7	93	43	8	43	57	74	57	86	42	60	96	90	59	17
56	49	66	15	17	74	51	89	4	74	26	75	27	99	79	72
30	76	8	54	80	65	97	50	14	68	86	33	19	51	63	19
20	52	49	51	87	72	49	77	71	31	51	71	92	9	76	77
13	51	93	88	19	14	15	40	57	60	54	69	76	31	51	29
39	5	52	55	18	94	2	40	8	61	49	44	35	19	57	80
99	52	1	2	70	66	10	24	66	68	0	37	74	2	12	87
7	46	41	9	63	59	56	62	3	20	56	21	3	25	46	48
18	76	81	96	91	65	38	83	10	89	2	54	21	43	4	53

# Deep Neural Networks




# Deep Neural Networks






# Deep Neural Networks



# Deep Neural Networks



# Deep Neural Networks

- ▶ We want to classify images into one of  $k$  classes

# Deep Neural Networks

- ▶ We want to classify images into one of  $k$  classes
- ▶ Extract hierarchical features


# Deep Neural Networks

- ▶ We want to classify images into one of  $k$  classes
- ▶ Extract hierarchical features

Edges


# Deep Neural Networks

- ▶ We want to classify images into one of  $k$  classes
- ▶ Extract hierarchical features




# Deep Neural Networks

- ▶ We want to classify images into one of  $k$  classes
- ▶ Extract hierarchical features



# Deep Neural Networks

- ▶ We want to classify images into one of  $k$  classes
- ▶ Extract hierarchical features





Why reinvent the wheel?

# Deep Neural Networks

Why reinvent the wheel?

13	81	43
60	83	34
32	46	99


# Deep Neural Networks

Why reinvent the wheel?

77
34
43
38
58
9
40
76
56


# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?



# Deep Neural Networks


Why reinvent the wheel?






# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?



# Deep Neural Networks


Why reinvent the wheel?






# Deep Neural Networks

Why reinvent the wheel?




# Deep Neural Networks

Why reinvent the wheel?



# Deep Neural Networks

Why reinvent the wheel?




Time for ze kitteh



This is a cat ♥

Time for ze kitteh



Still a cat ♥♥

Time for ze kitteh



Half cat / half salad ♥♥♥♥

Time for ze kitteh




Many cats ♥♥♥♥♥

## Weight Sharing




# Convolutional Neural Networks

## Weight Sharing




# Convolutional Neural Networks

## Weight Sharing




# Convolutional Neural Networks

## Weight Sharing




# Convolutional Neural Networks

## Weight Sharing




# Convolutional Neural Networks

## Weight Sharing



# Convolutional Neural Networks: Mechanics



Image

# Convolutional Neural Networks: Mechanics

0.4	0.9	0.1
0.7	0.2	0.6
0.8	0.3	0.5

Image

# Convolutional Neural Networks: Mechanics

0.4	0.9	0.1
0.7	0.2	0.6
0.8	0.3	0.5

Image

5.7	2.4
3.1	0.9

Kernel



# Convolutional Neural Networks: Mechanics

$5.7 \cdot 0.4$	$2.4 \cdot 0.9$	0.1
$3.1 \cdot 0.7$	$0.9 \cdot 0.2$	0.6
0.8	0.3	0.5

Image

# Convolutional Neural Networks: Mechanics

$5.7 \cdot 0.4$	$2.4 \cdot 0.9$	0.1
$3.1 \cdot 0.7$	$0.9 \cdot 0.2$	0.6
0.8	0.3	0.5

Image

6.79

Output

# Convolutional Neural Networks: Mechanics

0.4	$5.7 \cdot 0.9$	$2.4 \cdot 0.1$
0.7	$3.1 \cdot 0.2$	$0.9 \cdot 0.6$
0.8	0.3	0.5

Image

6.79	6.53
------	------

Output

# Convolutional Neural Networks: Mechanics

0.4	0.9	0.1
$5.7 \cdot 0.7$	$2.4 \cdot 0.2$	0.6
$3.1 \cdot 0.8$	$0.9 \cdot 0.3$	0.5

Image

6.79	6.53
7.67	

Output

# Convolutional Neural Networks: Mechanics


0.4	0.9	0.1
0.7	$5.7 \cdot 0.2$	$2.4 \cdot 0.6$
0.8	$3.1 \cdot 0.3$	$0.9 \cdot 0.5$

Image


6.79	6.53
7.67	3.96

Output


# Convolutional Neural Networks: Mechanics



# Convolutional Neural Networks: Mechanics




# Convolutional Neural Networks: Mechanics







# Convolutional Neural Networks: Mechanics




# Convolutional Neural Networks: Mechanics




# Convolutional Neural Networks: Mechanics




# Convolutional Neural Networks: Mechanics




# Convolutional Neural Networks: Mechanics




# Convolutional Neural Networks: Mechanics



# Convolutional Neural Networks: Mechanics




# Convolutional Neural Networks: Mechanics







# Convolutional Neural Networks: Mechanics



# Convolutional Neural Networks: Mechanics



# Convolutional Neural Networks: Mechanics

## Recipe for a Convolutional Layer

# Convolutional Neural Networks: Mechanics

## Recipe for a Convolutional Layer

- ▶ Ingredients

# Convolutional Neural Networks: Mechanics

## Recipe for a Convolutional Layer

- ▶ Ingredients

1. Image  $I$  with dimensions  $w \times h \times d$

## Recipe for a Convolutional Layer

► Ingredients

1. Image  $I$  with dimensions  $w \times h \times d$
2. A kernel (filter)  $K$  of size  $k \times l \times m$

## Recipe for a Convolutional Layer

- ▶ Ingredients

1. Image  $I$  with dimensions  $w \times h \times d$
2. A kernel (filter)  $K$  of size  $k \times l \times m$

- ▶ Cooking

## Recipe for a Convolutional Layer

- ▶ Ingredients
  1. Image  $I$  with dimensions  $w \times h \times d$
  2. A kernel (filter)  $K$  of size  $k \times l \times m$
- ▶ Cooking
  - ▶ Put the image into the oven at  $150^{\circ}\text{C}$



## Recipe for a Convolutional Layer

- ▶ Ingredients
  1. Image  $I$  with dimensions  $w \times h \times d$
  2. A kernel (filter)  $K$  of size  $k \times l \times m$
- ▶ Cooking
  - ▶ Don't put the image into the oven at 150°C

## Recipe for a Convolutional Layer

- ▶ Ingredients

1. Image  $I$  with dimensions  $w \times h \times d$
2. A kernel (filter)  $K$  of size  $k \times l \times m$

- ▶ Cooking

- ▶ Don't put the image into the oven at  $150^{\circ}\text{C}$
- ▶ Slide the kernel across the image

## Recipe for a Convolutional Layer

- ▶ Ingredients

1. Image  $I$  with dimensions  $w \times h \times d$
2. A kernel (filter)  $K$  of size  $k \times l \times m$

- ▶ Cooking

- ▶ Don't put the image into the oven at  $150^{\circ}\text{C}$
- ▶ Slide the kernel across the image
- ▶ Compute the “dot product” for each configuration

# Convolutional Neural Networks: Pooling

# Convolutional Neural Networks: Pooling

66	2
6	32

# Convolutional Neural Networks: Pooling

66	2
6	32

# Convolutional Neural Networks: Pooling

50	17
66	2

# Convolutional Neural Networks: Pooling

9	50
33	66



# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

## Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

66
----

# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

66	79
----	----

# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

66	79
66	

# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

66	79
66	128

# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

66	79
66	128

- *Pooling* achieves translational invariance

# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

66	79
66	128

- ▶ *Pooling* achieves translational invariance
- ▶ A form of downsampling



# Convolutional Neural Networks: Pooling

5	19	69
66	2	79
6	32	128

66	79
66	128

- ▶ *Pooling* achieves translational invariance
- ▶ A form of downsampling
- ▶ Other pooling functions possible

# Convolutional Neural Networks: Architecture

INPUT  $\rightarrow$  [CONV+  $\rightarrow$  POOL?]+  $\rightarrow$  FC+  $\rightarrow$  OUTPUT



TensorFlow



- ▶ An open source deep learning library

# TensorFlow




- ▶ An open source deep learning library
- ▶ Released by Google in November 2015



- ▶ An open source deep learning library
- ▶ Released by Google in November 2015
- ▶ Especially suited to:  
“Large-scale machine learning on heterogeneous distributed systems”

# Computational Paradigms

# Computational Paradigms




**Computational Graphs**

$$\hat{y} = \tanh(\mathbf{x}^\top \mathbf{w} + b)$$



# Computational Paradigms




$$\hat{y} = \tanh(\mathbf{x}^\top \mathbf{w} + b)$$

## Computational Graphs

### 1. Operations

# Computational Paradigms




$$\hat{y} = \tanh(\mathbf{x}^\top \mathbf{w} + b)$$

## Computational Graphs

1. Operations
2. Tensors

# Computational Paradigms




$$\hat{y} = \tanh(\mathbf{x}^\top \mathbf{w} + b)$$

## Computational Graphs

1. Operations
2. Tensors
3. Variables

# Computational Paradigms




## Computational Graphs

1. Operations
2. Tensors
3. Variables
4. Sessions

$$\hat{y} = \text{session.run}(\tanh(\mathbf{x}^T \mathbf{w} + b))$$


# Execution Model

# Execution Model




## 1. Client

# Execution Model



1. Client
2. Master


# Execution Model



1. Client
2. Master
3. Workers




# Execution Model




1. Client
2. Master
3. Workers
4. Devices

# Execution Model



1. Client
2. Master
3. Workers
4. Devices

# Execution Model



1. Client
2. Master
3. Workers
4. Devices

# Visualization Tools

# Visualization Tools

- ▶ Deep Neural Networks have the tendency of being . . . deep

# Visualization Tools


- ▶ Deep Neural Networks have the tendency of being . . . deep
- ▶ Easy to drown in the complexity of an architecture


# Visualization Tools


- ▶ Deep Neural Networks have the tendency of being . . . deep
- ▶ Easy to drown in the complexity of an architecture
- ▶ > 36,000 nodes for Google's *Inception* model














# TensorBoard to the Rescue

# Walkthrough

## LeNet5

INPUT  $\rightarrow$  [CONV  $\rightarrow$  POOL] $\{2\}$   $\rightarrow$  FC  $\rightarrow$  OUTPUT

How do I continue?



# Resources

# Resources

- ▶ MOOCs
  - ▶ Machine Learning by Andrew Ng @ Coursera
  - ▶ Deep Learning by Google @ Udacity
  - ▶ Machine Learning Nanodegree @ Udacity

# Resources

## ▶ MOOCs

- ▶ Machine Learning by Andrew Ng @ Coursera
- ▶ Deep Learning by Google @ Udacity
- ▶ Machine Learning Nanodegree @ Udacity

## ▶ Websites

- ▶ <http://colah.github.io>
- ▶ <http://cs231n.github.io>
- ▶ <http://karpathy.github.io>
- ▶ <http://www.deeplearningbook.org>
- ▶ <https://www.kaggle.com>
- ▶ <https://www.tensorflow.org>

# PyCon Germany

- ▶ 29-30 October in Munich
- ▶ Talks on machine learning, deep learning and data science
- ▶ 15% off: VISITMUC16

# Stay in Touch!

- ▶ `peter@goldsborough.me`
- ▶ `linkedin.com/in/petergoldsborough`
- ▶ `github.com/goldsborough`
- ▶ `@peterawks`

# Stay in Touch!

- ▶ `peter@goldsborough.me`
- ▶ `linkedin.com/in/petergoldsborough`
- ▶ `github.com/goldsborough`
- ▶ `@peterawks`

`github.com/peter-can-talk/pycon-uk-2016`

# Q & A